Goal

Update Firewall Rules to support AKS Cluster for inbound and outbound connections

Note: Run the lab from the same directory where you found this instructions file

Steps

Create Log Analytics for Azure Firewall

1. Lets set up diagnostics and import the dashboard file as described here: https://docs.microsoft.com/en-us/azure/firewall/tutorial-diagnostics. The goal here is to be able to analyze what Azure FW is seeing from both an ingress and egress perspective throughout the lab.

```
#Setup AzFW LogAnalytics
#https://docs.microsoft.com/bs-latn-ba/azure/firewall/log-analytics-samples
```

Go to Azure Firewall blade and click Diagnostics setttings. Click Turn on Diagnostics:

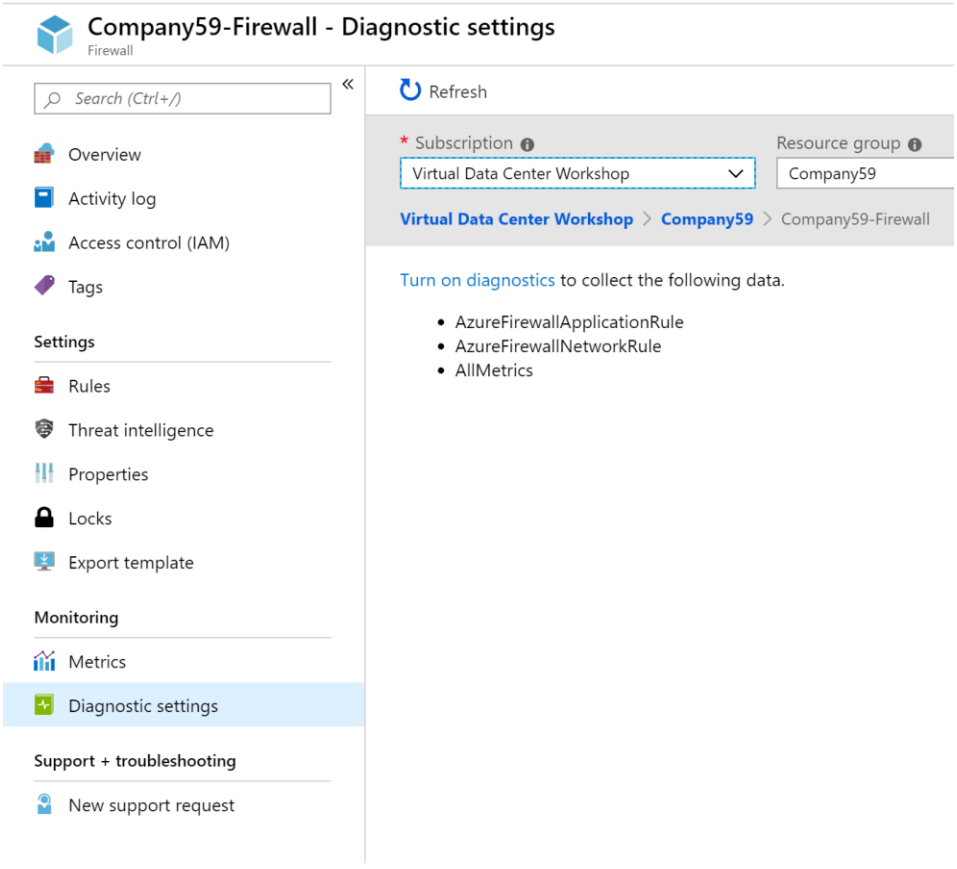

2. Fill out a name for the Diagnostic settings, including click Send to Log Analytics.

Make sure your subscription and Log Analytics Workspace is selected.

Finally, click on AzureFirewallApplicationRule, AzureFirewallNetworkRule and All Metrics.

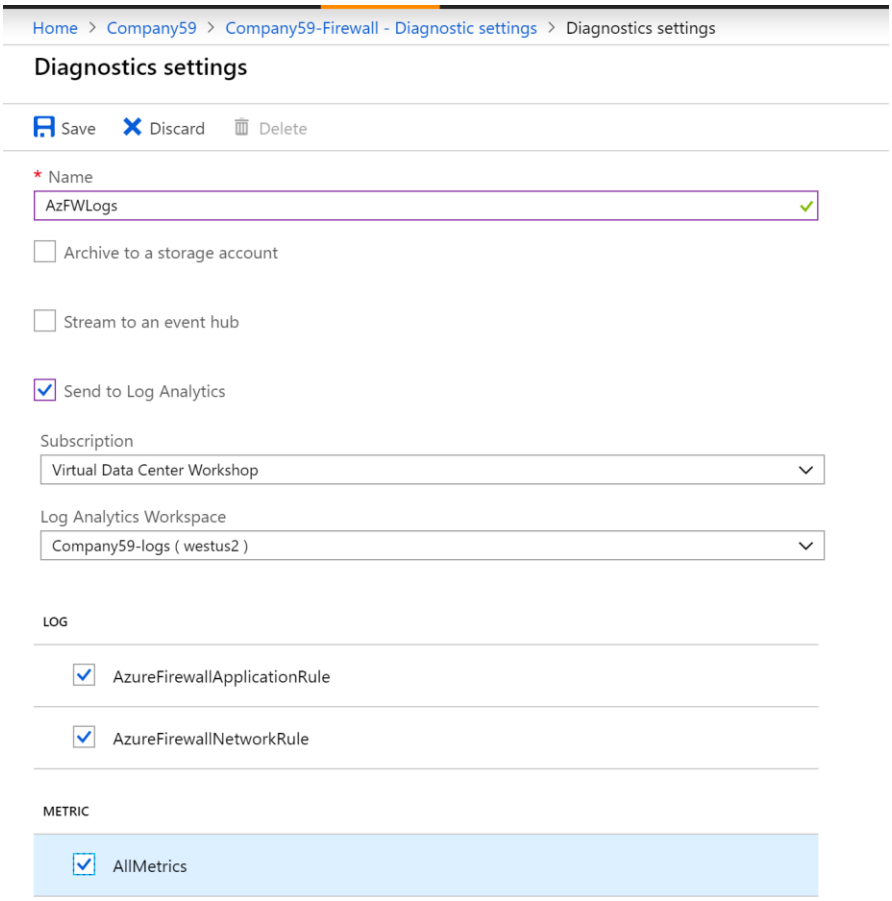

3. Navigatge to LogAnalytics Workspace blade. Click on View Designer. Click import and select file AzureFirewall.omsview.
Remember to click save and also save as Dashboard.

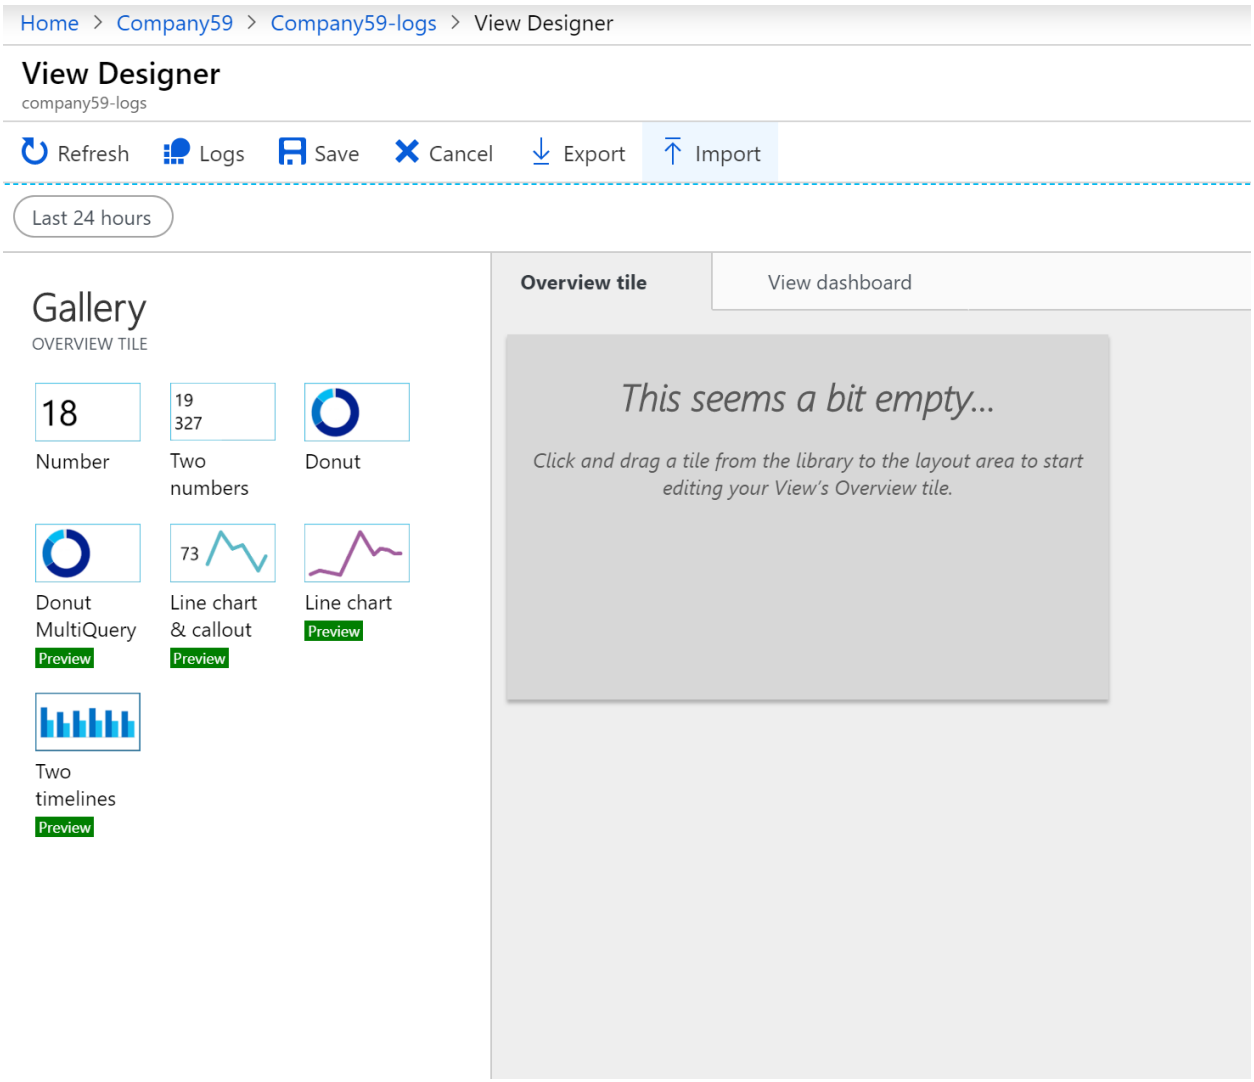

4. Click save and then click the pin to save to Dashboard

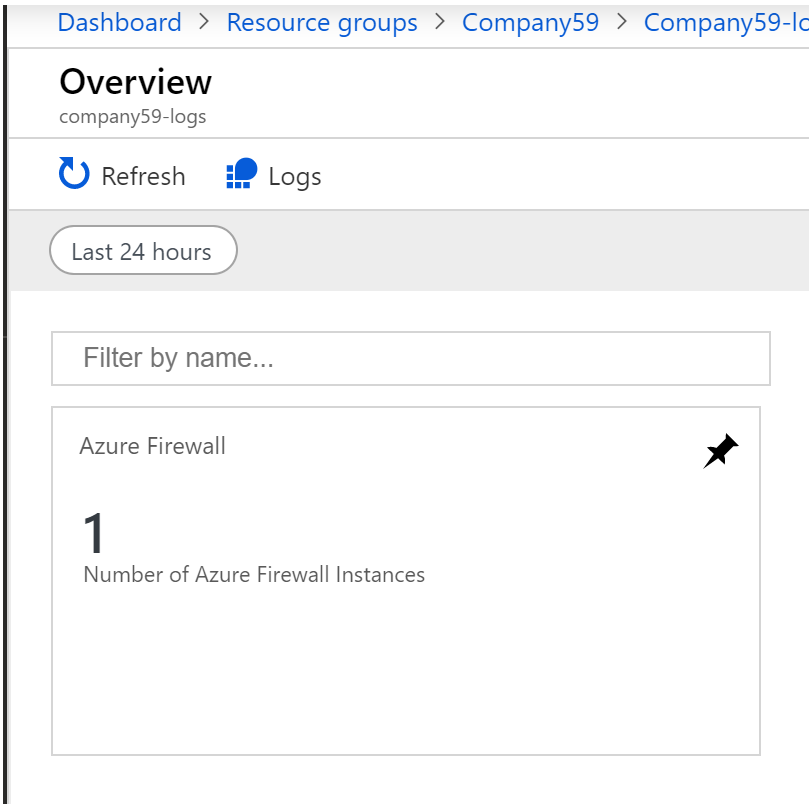

Create Variables and Provision Azure Firewall Rules and Associate to AKS Subnet.

The end result will look like the below diagram. The lab will require some steps to configure the vnet, subnets, routetable, firewall rules and azure kubernetes services which are described below and can be adapted to any kubernetes installation on azure:

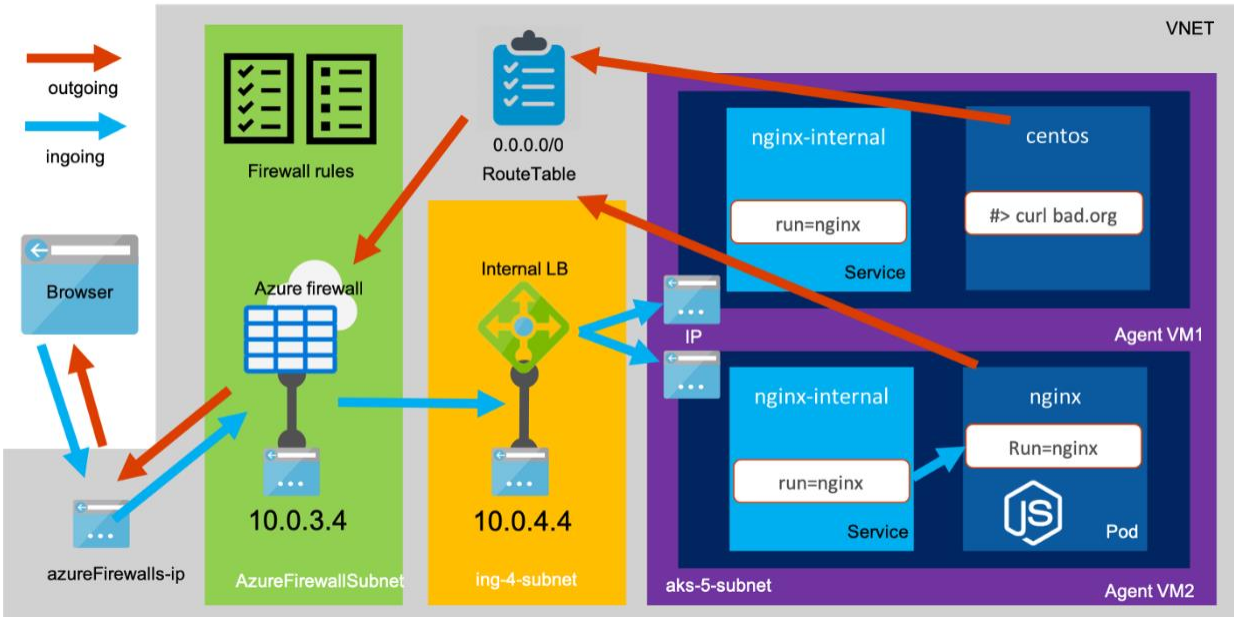

We will create variables to use throughout the remaining lab. Remember to update with your specific information replacing the XX with your ID:

1. Update information and copy and paste into your terminal:

```
$RG="CompanyXX" # here enter the resources group name of your AKS cluster

$LOC="westus"   # here enter the datacenter location

$NAME="CXX-Spoke-Cluster" # here enter the name of your kubernetes resource

$VNET_NAME="CXX-Spoke-VNet" # here enter the name of your vnet

# DO NOT CHANGE FWSUBNET_NAME - This is currently a requirement for Azure Firewall.

$FWSUBNET_NAME="AzureFirewallSubnet"

$FWNAME="CompanyXX-Firewall"

$FWPUBLICIP_NAME="CompanyXX-Firewall-pip"

$AKSSUBNET_NAME="Cluster"

$FWROUTE_TABLE_NAME="CXX-VNet-rt-fw"
```

1. Create variable to grab existing Azure firewall public and private IP address:

```
$FWPUBLIC_IP=(az network public-ip show -g $RG -n $FWPUBLICIP_NAME --query "ipAddress" -o tsv)

$FWPRIVATE_IP=(az network firewall show -g $RG -n $FWNAME --query
"ipConfigurations[0].privateIpAddress" -o tsv)
```

2. Validate Azure Firewall IP Address Values - This is more for awareness so you can help connect the networking dots

```
echo $FWPUBLIC_IP

echo $FWPRIVATE_IP
```

3. Create rules on Azure Firewall to allow outbound connections to support AKS control plane:

```
# There is a private preview feature that allows you to whitelist what IPs are allowed to access the Master API server. Secondly, this summer Azure will preview private-link to provide a single customer owned private IP address to identify Master API server which will further secure control plane traffic.

# Add the Azure Firewall extension to Azure CLI in case you do not already have it.
```

az extension add --name azure-firewall

4. Create the Outbound Network Rule from Worker Nodes to Control Plane

az network firewall network-rule create -g $RG -f $FWNAME --collection-name 'aksfwnr' --name "allow network" --protocols 'TCP' --source-addresses '*' --destination-addresses '*' --destination-ports 22 443 --action allow --priority 100

5. Add Application FW Rules for Egress Traffic that supports AKS control plane

az network firewall application-rule create --firewall-name $FWNAME --collection-name "aksbasics" --name "allow network" --protocols http=80 https=443 --source-addresses "*" --resource-group $RG --action "Allow" --target-fqdns "*eastus.azmk8s.io" "*auth.docker.io" "*cloudflare.docker.io" "*registry-1.docker.io" "k8s.gcr.io" "storage.googleapis.com" "*cloudflare.docker.com" --priority 100

az network firewall application-rule create --firewall-name $FWNAME --collection-name "aksextended" --name "allow network" --protocols http=80 https=443 --source-addresses "*" --resource-group $RG --action "Allow" --target-fqdns "login.microsoftonline.com" "*.management.azure.com" "mcr.microsoft.com" "download.opensuse.org" "*.azureedge.net" "*.ubuntu.com" "*azurecr.io" "*blob.core.windows.net" --priority 101

6. Associate your AKS subnet with a route table to route 0.0.0.0/32 to Azure Firewall.

az network vnet subnet update -g $RG --vnet-name $VNET_NAME --name $AKSSUBNET_NAME --route-table $FWROUTE_TABLE_NAME

Now all outgoing traffic will be filtered and you can check that by launching a pod and see if you can curl the outside internet.

Deploy AKS Ingress Controller with Azure Load-Balancer

1. Enter the following AzCLI to get status on your cluster and current nodes deployed. Check Provisioning Status of AKS Cluster - ProvisioningState should say 'Succeeded'

az aks list -o table

2. Get AKS Credentials so kubectl works

az aks get-credentials -g $RG -n $NAME --admin

3. Setup alias to kubectl becuase I am lazy

alias k="kubectl"

4. Get Nodes

k get nodes -o wide

5. Let's check if the Azure Firewall is properly working for egress traffic from a POD. Check Egress Traffic through Firewall

kubectl apply -f TestFW.yml

kubectl get po -o wide

kubectl exec -it centos -- /bin/bash

This curl command should work

curl www.ubuntu.com

This curl command should fail showing that not even Superman can get through Azure Firewall

curl www.espn.com

exit

View Azure Firewall LogAnalytics Dashboard.

1. Navigate to Dashboard and click on your saved Dashboard AzureFirewall. View what urls are being accessed and what IPs are being used.

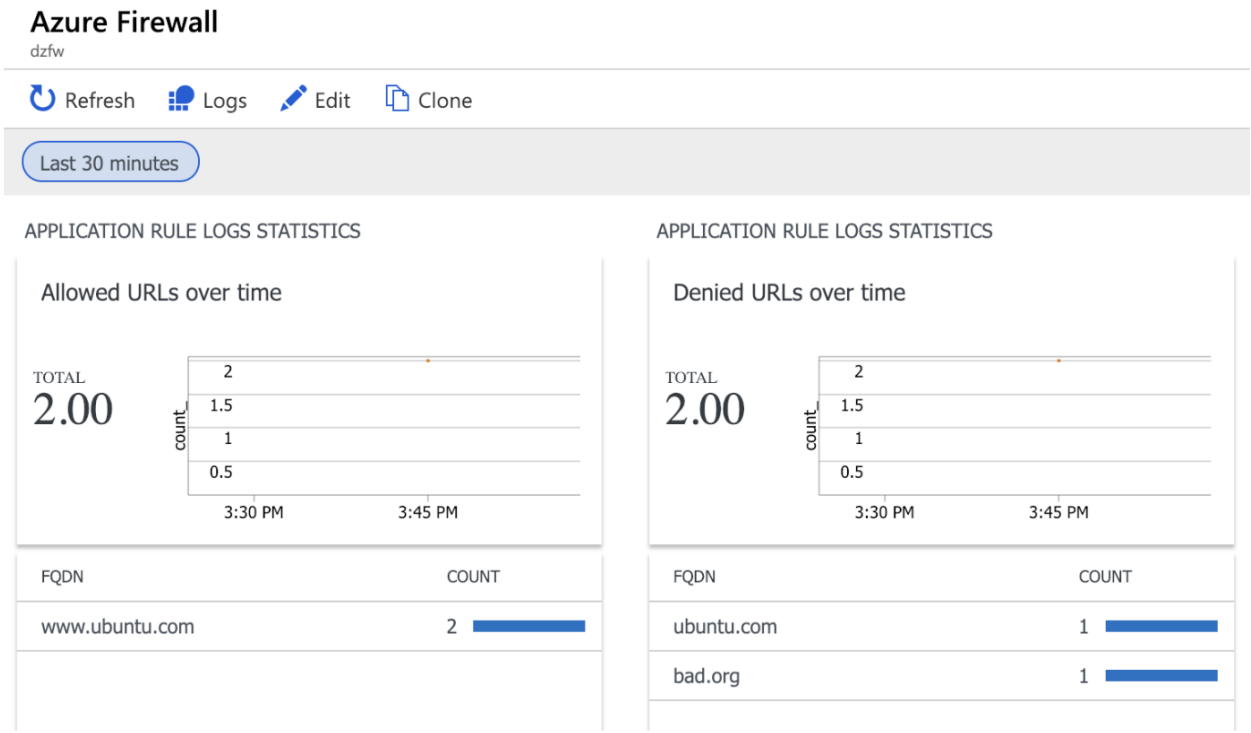

Deploy a new Service that is securely accessible via the Internet

A little about this workload. The first service is a web front-end written in .NET Core that sets up a SignalR Hub for receiving real-time communication from the back-end worker node, which in turn relays that to the web client.

The second service is a tensorflow model that classifies pictures of fruit that are located in Azure Files. Each time an image has been classified it communicates back to the SignalR Hub setup on the web front-end which then relays that to the client.

This setup is different than your typical web to database communication as it flows the other way from back-end to front-end. Keep those communication paths in mind when we get to updating variables for communication between servics as well as service segmentation.

1. First lets make sure we have proper DNS setup. Get the Public IP address of your Azure Firewall

```
$IP=(az network public-ip show -g $RG -n $FWPUBLICIP_NAME --query "ipAddress" -o tsv)

Echo $IP
```

2. Name to associate with public IP address of your new service you are about to

```
$DNSNAME="aksazfw"
```

3. Navigate to the Public IP of your Azure Firewall

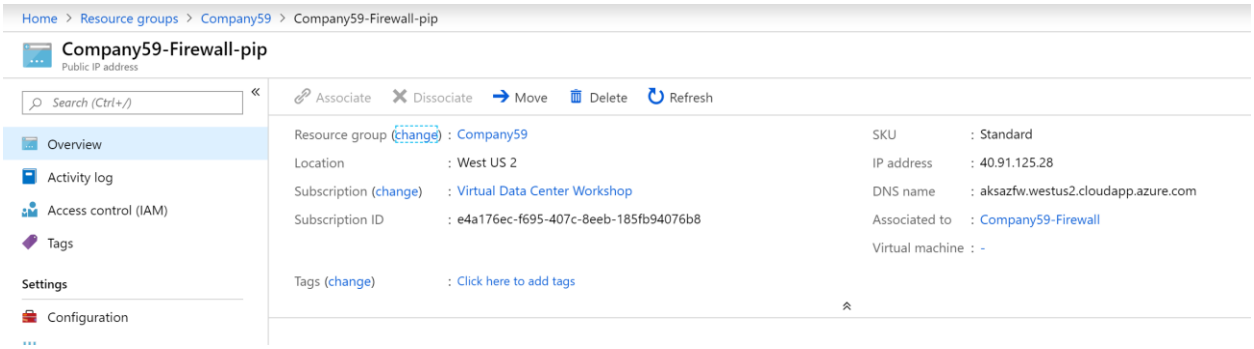

4. Click configuration and update with your dns

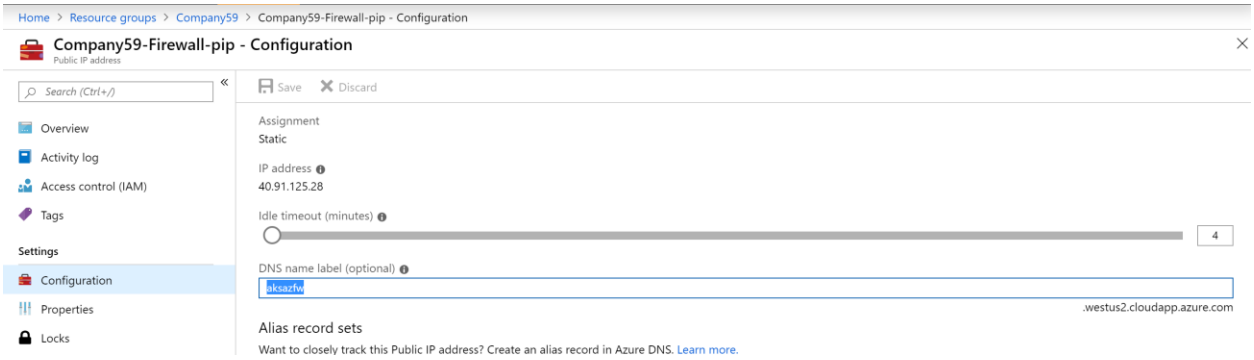

5. Next before deploying the Pods we need to update winter-ready-web2.yaml We need to edit line 45 and enter an available private IP address from Cluster subnet. #Make sure subnet name and address space is updated in yaml file.

```
32        dnsPolicy: ClusterFirst
33    ---
34    apiVersion: v1
35    kind: Service
36    metadata:
37      name: imageclassifierweb
38      labels:
39        app: imageclassifierweb
40      annotations:
41        service.beta.kubernetes.io/azure-load-balancer-internal: "true"
42        service.beta.kubernetes.io/azure-load-balancer-internal-subnet: "Cluster"
43    spec:
44      type: LoadBalancer
45      loadBalancerIP: "10.17.59.231"
46      ports:
47        - port: 80
48          targetPort: http
49          protocol: TCP
50          name: http
51      selector:
52        app: imageclassifierweb
53
```

6. Open up winterready-ingress-web2.yaml We need to edit the host entry (line 10 and 17) with the DNS name you created above.

```yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: winterready
  annotations:
    kubernetes.io/ingress.class: nginx
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - host: aksazfw.westus2.cloudapp.azure.com
    http:
      paths:
      - path: /
        backend:
          serviceName: imageclassifierweb
          servicePort: 80
  - host: aksaz.internal.cloudapp.azure.com
    http:
      paths:
      - path: /
        backend:
          serviceName: imageclassifierweb
          servicePort: 80

```

7. Now lets deploy the service. We are going to add Web Front-End first.

k apply -f winterready-web2.yaml

8. Next we are going to deploy the Ingress Controller which creates an NGINX service for L7 proxy as well as created an Azure internal Load-balancer.

k apply -f winterready-ingress-web.yaml

9. We need to upload the images in the fruit directory into an Azure Storage Account that is setup with a new Fileshare called Fruit where we store the images.

   a. Click on your storage account and then navigate to firewall and virtual networks. Click all networks and save.

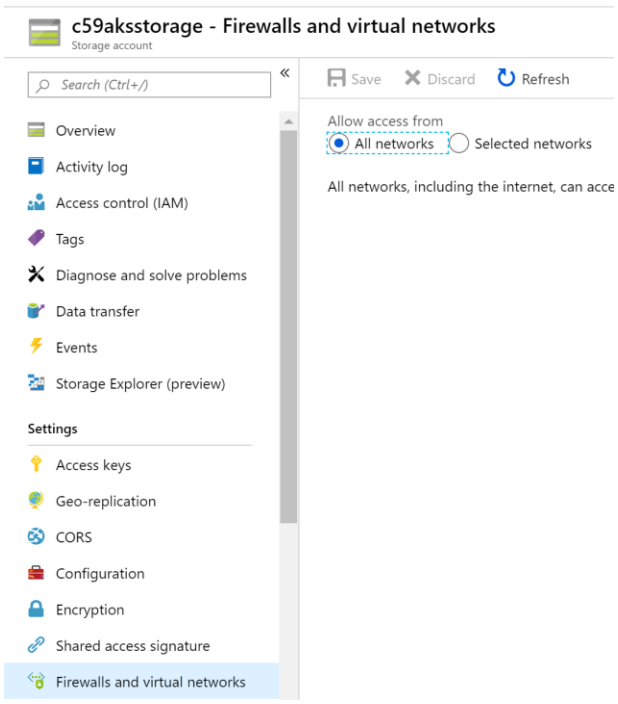

a. Click Overview and then click on Files. Click +File share and enter the name fruit with 1GB for Quota.

b. Click the new fileshare Fruit and click upload. Select all fruit images and upload.

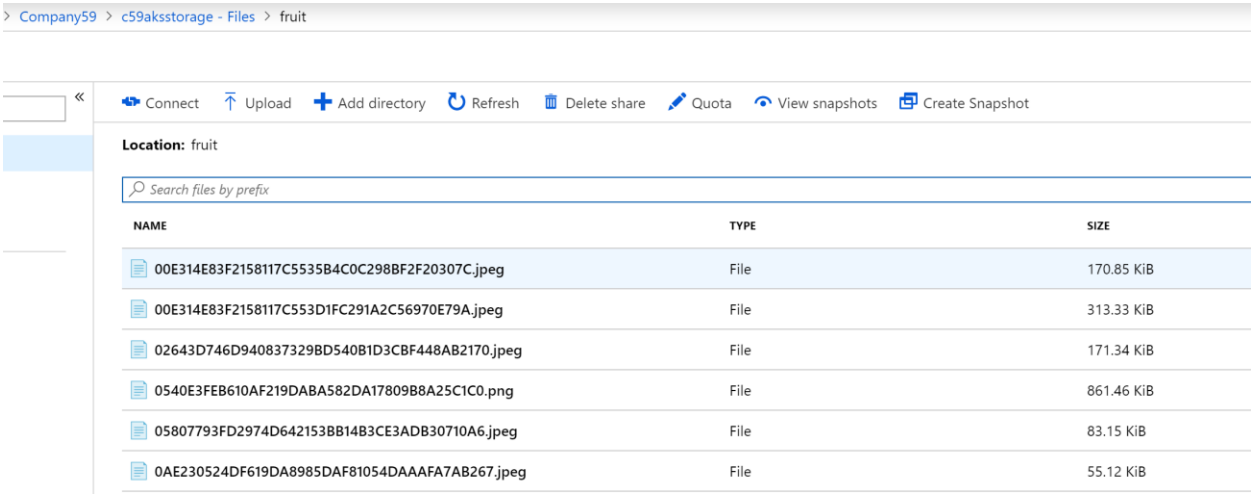

10. We need to secure access to the file share. We will add in the storage account and keys to AKS.

Update <STORAGE_ACCOUNT_NAME_GOES_HERE> and <STORAGE_ACCOUNT_KEY_GOES_HERE> below with the correct credentials from the Azure Storage account created above. Make sure you update the AzCLI with your storage acct name and key. Storage acct key can be found here:

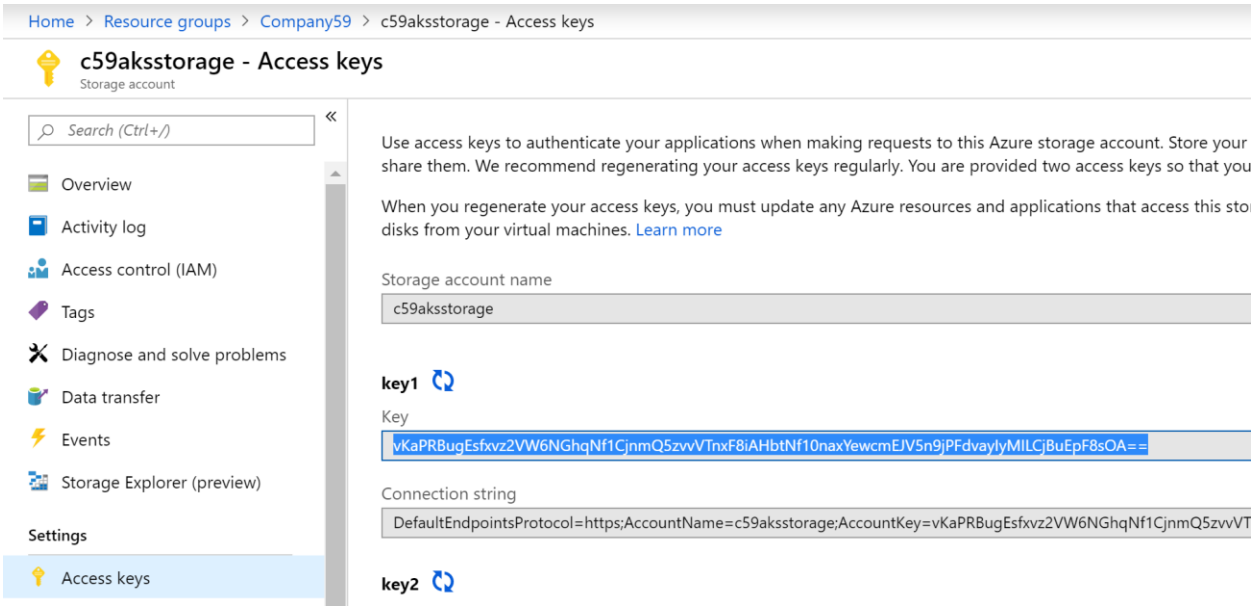

kubectl create secret generic fruit-secret --from-literal=azurestorageaccountname=c59aksstorage --from-literal=azurestorageaccountkey=vKaPRBugEsfxvz2VW6NGhqNf1CjnmQ5zvvVTnxF8iAHbtNf10naxYewcmEJV5n9jPFdvaylyMILCjBuEpF8sOA==

11. Time to deploy the Worker Back-End:

kubectl apply -f .\winterready-worker.yaml

12. Lets take a look at the services we created. Pay close attention to the service with the "external-ip" which defines the private IP on Azure LB. You can also see the hosts FQDNs and Secret key for accessing Azure Files.

kubectl get po,svc,ingress,deploy,secrets

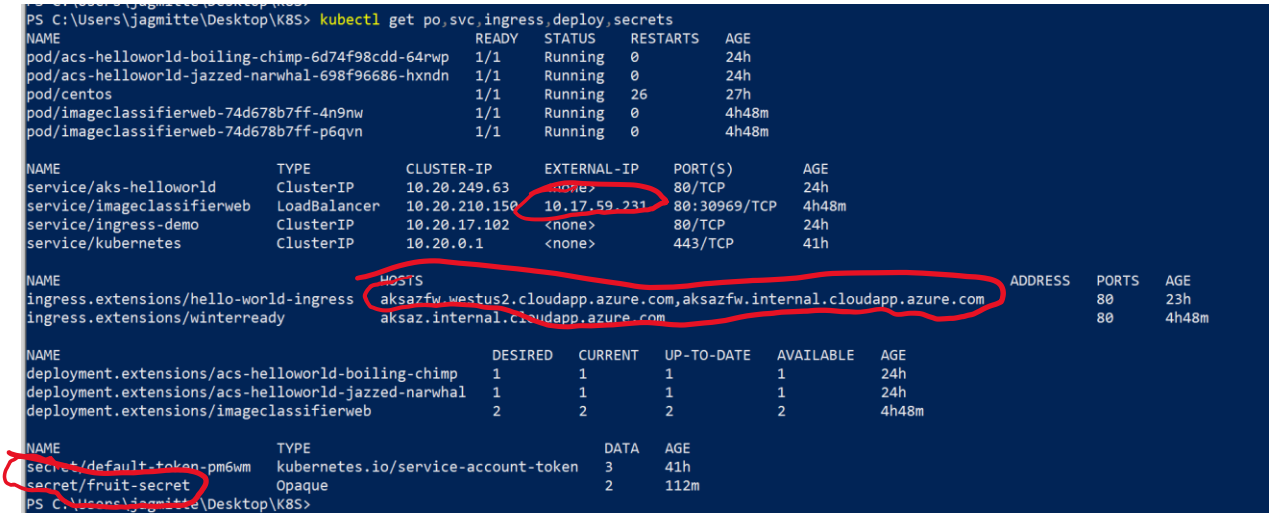

13. Now we can retrieve the internal load balancer ip and register it in the azure firewall as a Dnat rule.

```
$SERVICE_IP=$(kubectl get svc imageclassifierweb --template="{{range
.status.loadBalancer.ingress}}{{.ip}}{{end}}")
```

14. Lets up the Azure FW rule to permit traffic from the Internet with DNAT.

```
az network firewall nat-rule create  --firewall-name $FWNAME --collection-name "inboundlbrules" --
name "allow inbound port 80" --protocols "TCP" --source-addresses "*" --resource-group $RG --action
"Dnat"  --destination-addresses $IP --destination-ports 80 --translated-address $SERVICE_IP --translated-
port "80"  --priority 600
```

15. Now you can access the internal service by going to the DNS assigned to the public IP of your azure firewall on port 80
http://$DNSNAME.westus2.cloudapp.azure.com

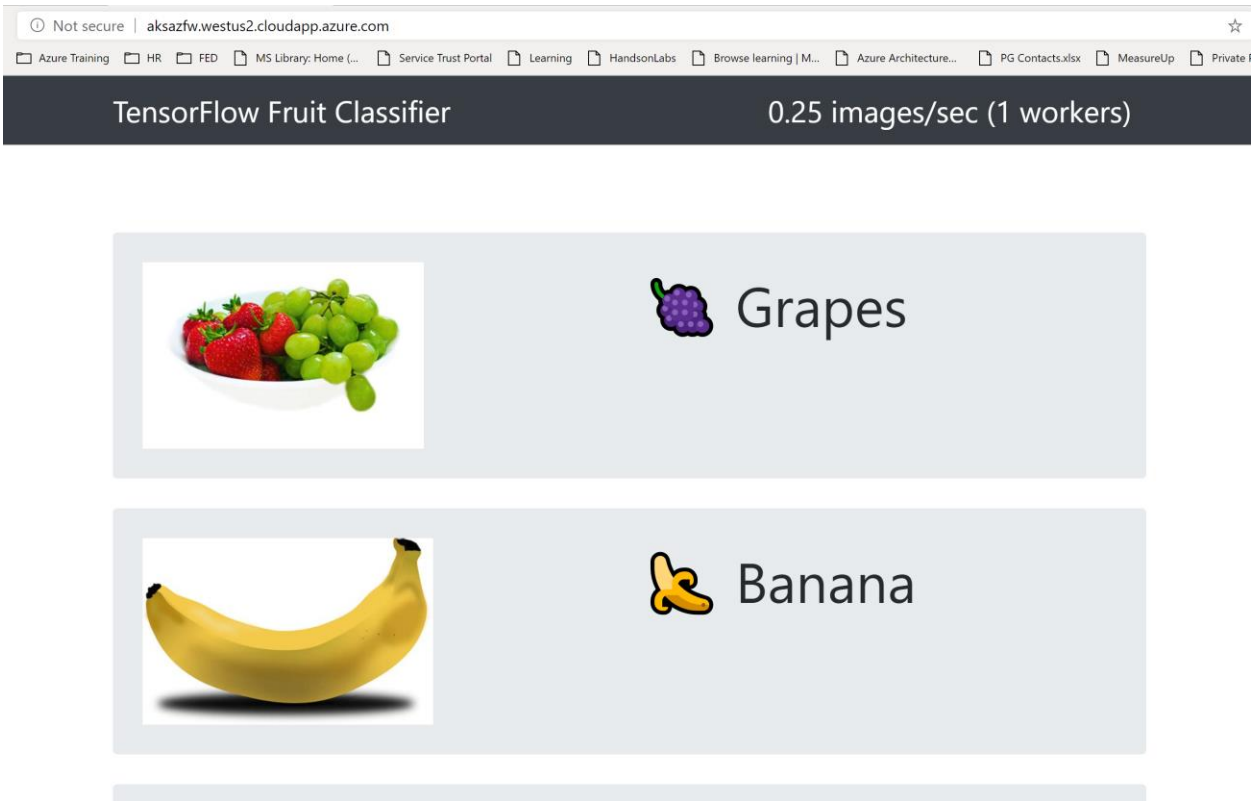

Last item is accessing the service via ExpressRoute Private-Peering.

1. Navigate to VM that is on-premise. Update the host file on the host and update with the private FQDN you created above with the private IP address of Azure Load Balancer.

2. Search for notepad and right click and open as administrator

3. Open the host file using the following path:

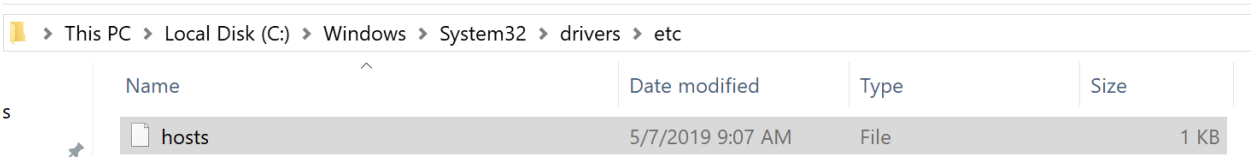

Edit the host file with the private IP address of your Azure LoadBalancer and Internal FQDN you defined earlier in the hello-world-ingress.yml

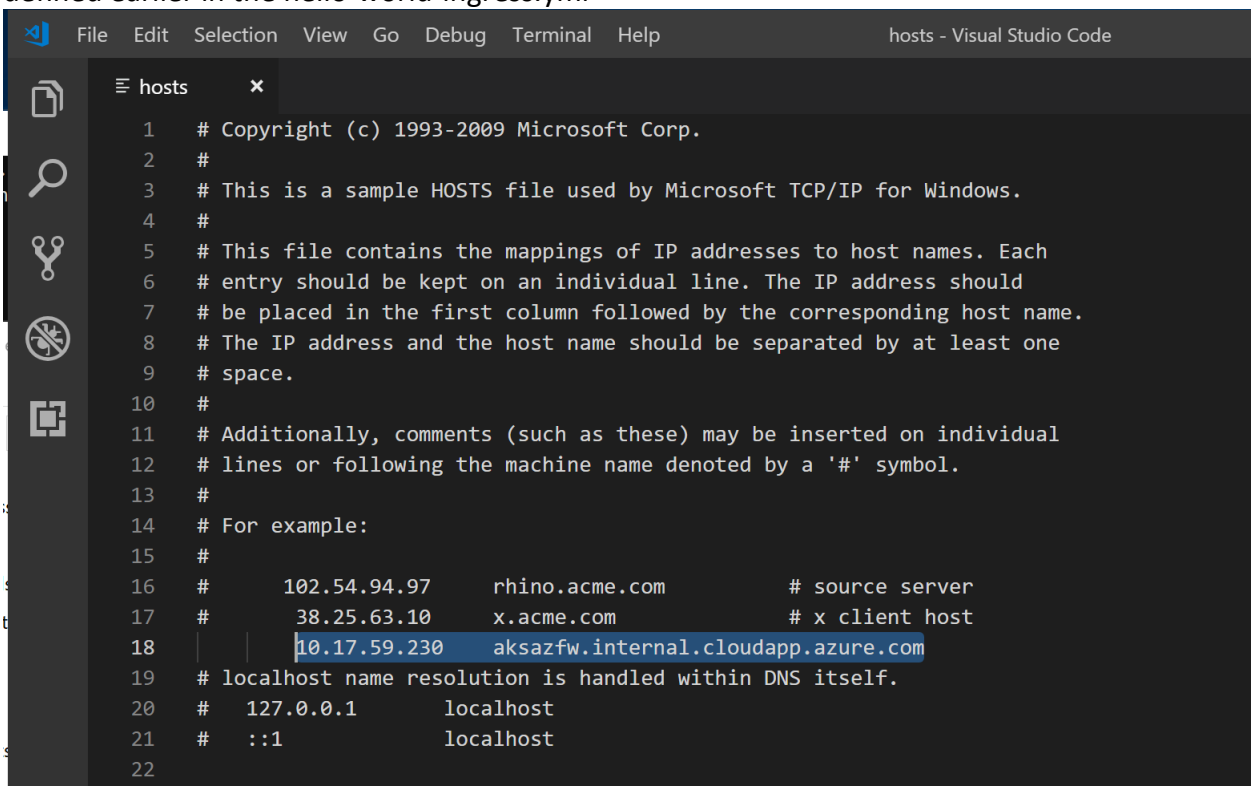

4. Create a new Azure Firewall rule that permits traffic to the Azure Load-Balancer IP address from any source.

az network firewall network-rule create -g $RG -f $FWNAME --collection-name 'onpremises' --name "allow network" --protocols 'TCP' --source-addresses '*' --destination-addresses $SERVICE_IP --destination-ports 80 --action allow --priority 300

5. Create a route-table that routes the private IP of the Azure Load Balancer with its next-hop being the private IP address of Azure Firewall.

az network route-table create --resource-group $RG --name OnPremisesAKS

az network route-table route create --name AKS --resource-group $RG --route-table-name OnPremisesAKS --address-prefix "$SERVICE_IP/32" --next-hop-type VirtualAppliance --next-hop-ip-address $FWPRIVATE_IP

6. Assign the route-table to the Gateway Subnet to service-chain Azure Firewall in path to Azure Load-Balancer.

az network vnet subnet update --vnet-name C59-VNet --name GatewaySubnet --resource-group $RG --route-table OnPremisesAKS

7. Open a browser on test vm and access the internal FQDN.

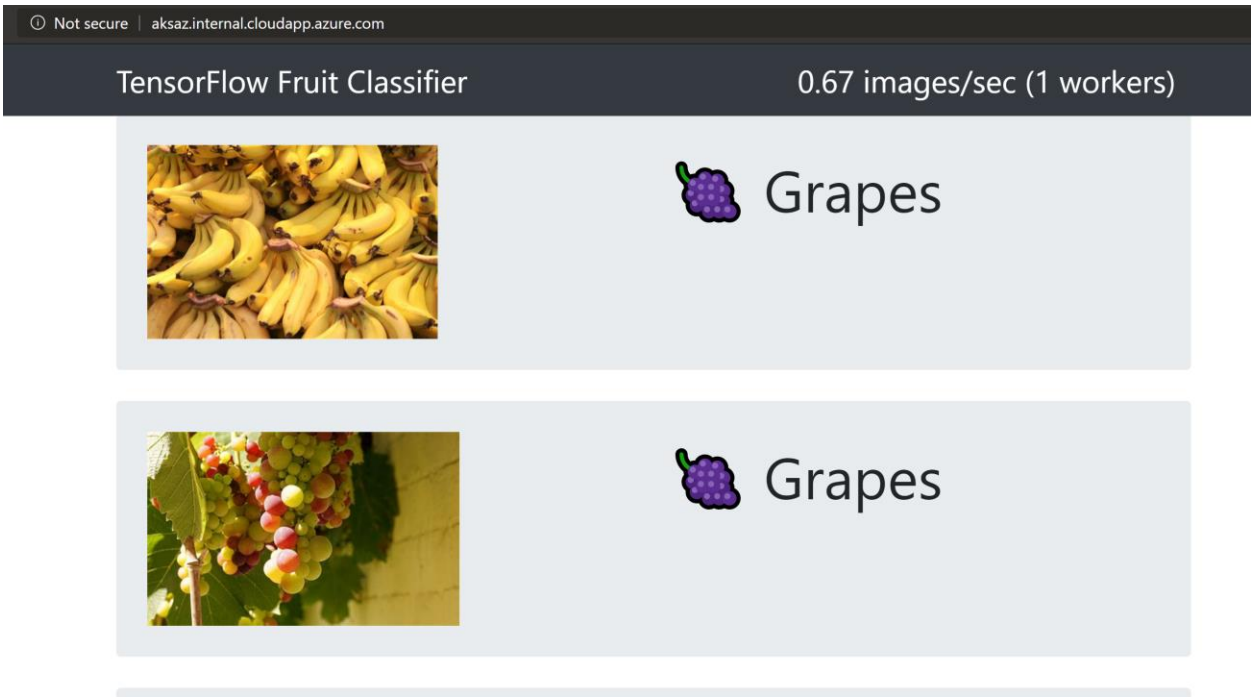

8. Review Logs within Azure Log Analytics to verify. Navigate to Dashboard and click on Azure FW. Scroll over to Network Rule Logs Statistics. Find the private IP or your on-premises VM:

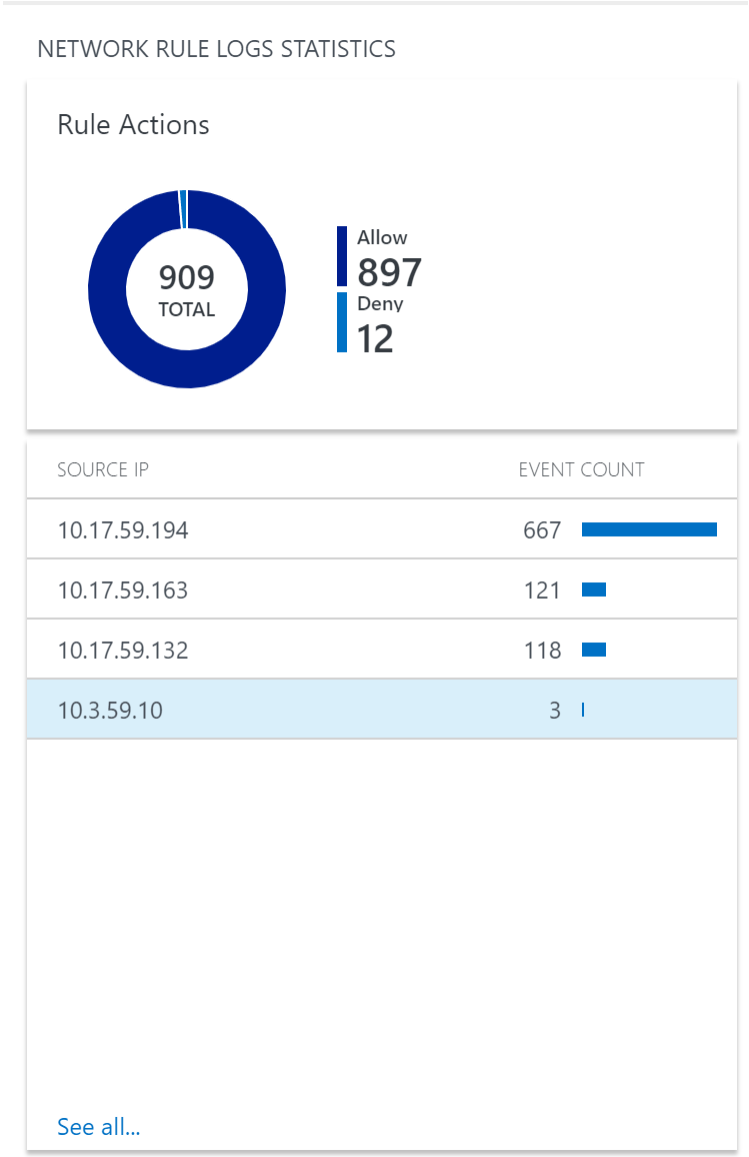

Review the status of all the containers in the deployment. Run the following commands:

kubectl get po,svc,ingress,deploy,secrets --all-namespaces

kubectl get pods -o wide
kubectl describe pod “PODNAME”

kubectl get pods -o wide --all-namespaces
kubectl get svc
kubectl describe svc “SVCNAME”

kubectl get ingress
kubectl get ingress hello-world-ingress
kubectl describe ingress hello-world-ingress